

Acceleration of Metropolis-Hasting Algorithm

Tairan

May 3, 2018

Motivation

Bayesian Framework

Let's begin with the brief introduction to the Bayesian framework. In Bayesian statistics, we would like to figure out what the parameter space will look like after considering the available data. From the mathematics point of view, we can summarize the Bayesian framework as:

$$f(\theta|X) = \frac{f(X|\theta) \times f(\theta)}{f(X)}$$

The notation can be explained as following:

- $f(\theta)$, or the prior, is our initial belief about the distribution of the parameter.
- $f(X|\theta)$, or the likelihood, is the distribution of the data given the parameters.
- $f(\theta|X)$, or the posterior, is the distribution of the parameter given the data

Since the denominator $f(X)$ doesn't contain the parameter θ , the equation can be simplified as:

$$\text{Posterior} \propto \text{Likelihood} \times \text{Prior}$$

However, the posterior distribution sometimes is hard to solve in mathematics form. Then, we need to introduce the sampling algorithm to explore its equilibrium distribution.

Sampling Algorithm

The Markov chain Monte Carlo (MCMC) methods include a class of algorithm for sampling from a probability distribution. We construct the Markov chain to depict the desired distribution, and then study its properties and conduct the statistics test. The most common applied sampling algorithms in MCMC are Gibbs sampling, Metropolis-Hasting sampling, and slice sampling. (Wikipedia contributors 2018b)

- Gibbs sampling
 - When to apply:
Gibbs sampling works well when the conditional distribution for each random variable is in an analytic form, like belonging to a known distribution. In this case, we can directly generate the sample distribution from this known distribution.
 - Strength:
If Gibbs sampling is applicable, it is definitely the most efficient sampling algorithm.
 - Weakness:
The requirement for the conditional distribution is extremely higher than other sampling algorithms.
- Metropolis-Hasting sampling
 - When to apply:
Metropolis-Hasting sampling can work in most cases.

- Strength:
You don't need to worry too much about the mathematics deductions for the conditional distribution.
- Weakness:
Building a well-tuned proposal distribution demands experience and might take long time. If the step size is too small random walk causes slow decorrelation. If the step size is too large there is great inefficiency due to a high rejection rate.
- Slice sampling
 - When to apply:
Slice sampling performs well in unimodal distributions and some steep multi-modals distributions.
 - Strength:
Compared to Metropolis-Hasting sampling, there is no need to manually tune the proposal function or proposal standard deviation.
 - Weakness:
When it comes to multidimensional distributions, we can sample each random variable using one-dimensional slice sampling, in a manner similar to Gibbs sampling. However, this approach severely suffer from complication introduced by dimensionality.

As we can see, Metropolis-Hasting is the most widely applicable sampling algorithm. However, the evaluation of the posterior and the proposal distribution is a significant run time bottleneck. Then we are seeking a method that can reduce the run time of Metropolis-Hasting.

Introduction of C++ to sampling algorithm

C++'s advantage in computation efficiency

C++ is a general-purpose programming language. It has imperative, object-oriented and generic programming features, while also providing facilities for low-level memory manipulation.

The reason why C++ is fast in computation can be summarized as following:

- Compact Data Structure:
C++ is intrinsically stingy with memory – a C++ struct has no memory overhead if there are no virtual functions. Smaller things run faster due to caching, and are also more scalable.
- Template Meta-programming:
Huge gains to be made with the kind of evaluation you can implement with expression templates and the kind of memory and indirection savings you achieve by avoiding virtual functions with the curiously recurring template pattern.

Therefore, the ideal procedure to analyze the posterior distribution is to sample in C++ for computation efficiency and study in R for mass statistics analysis packages.

RCpp

The RCpp package provides R functions as well as C++ classes which offer a seamless integration of R and C++. Many R data types and objects can be mapped back and forth to C++ equivalents which facilitates both writing of new code as well as easier integration of third-party libraries. (Eddelbuettel et al. 2011)

The core packages in RCpp includes: (Wikipedia contributors 2018a)

- Rcpp
 - Rcpp provides a powerful API on top of R, permitting direct interchange of rich R objects between R and C++
 - Rcpp sugar gives syntactic sugar such as vectorised C++ expression

- Rcpp modules provide easy extensibility using declarations and Rcpp attributes greatly facilitates code integration. (Rcpp 2018)
- RcppArmadillo
 - Rcpp connects R with the powerful Armadillo templated C++ library for linear algebra
 - Armadillo aims towards a good balance between speed and ease of use, and its syntax is deliberately similar to Matlab which makes it easy to port existing code
- RcppEigen
 - RcppEigen gives R access to the high-performance Eigen linear algebra library
 - Eigen is templated and highly optimized, which provides a wide variety of matrix methods, various decomposition and includes support for sparse matrices.
- RInside
 - The RInside package provides C++ classes that make it easier to embed R in C++ code by wrapping the existing R embedding API in an easy-to-use C++ class
 - Over a dozen basic example are included in the package, as well as more specialized example showing use of RInside with MPI for parallel computing, Qt for cross-platform GUI, Wt for web applications, as well as the Armadillo and Eigen templated C++ libraries for linear algebras

Stan

Stan is a probabilistic programming language for statistical inference written in C++. The Stan language is used to specify a (Bayesian) statistical model with an imperative program calculating the log probability density function.

Stan implements gradient-based Markov chain Monte Carlo (MCMC) algorithms for Bayesian inference, stochastic, gradient-based variation Bayesian methods for approximate Bayesian inference, and gradient-based optimization for penalized maximum likelihood estimation.

The package **RStan** in R can be employed to connect R and Stan. Therefore, the sample distribution is simulated in Stan and then transport back to R.

An Example: Logistic Model

Artificial Data

The artificial data is generated as following:

- The sample size n is 50
- X_1 : continuous variable, subject to a normal distribution with mean 0 and variance 1
- X_2 : continuous variable, subject to a normal distribution with mean 0.5 and variance 1
- X_3 : continuous variable based on the absolute value of X_1
- Y : binary outcome, generated based on:

$$\text{logit}(E(Y)) = 1 + 3 \times X_1 + 3 \times X_2 - 3 \times X_3$$

Bayesian Modeling

Based on the artificial data, we apply the logistic regression model:

$$\begin{aligned} \text{logit}(E(Y)) &= \text{logit}(p_i) \\ &= \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 \end{aligned}$$

Also provide the parameter with noninformative prior:

$$\vec{\beta} \sim N_4 \left(\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 100 & 0 & 0 & 0 \\ 0 & 100 & 0 & 0 \\ 0 & 0 & 100 & 0 \\ 0 & 0 & 0 & 100 \end{bmatrix} \right)$$

By easy mathematics deduction, we can find out the posterior distribution as:

$$\log f(\vec{\beta}|X, Y) \propto \sum \{y_i \times \log(p_i) + (1 - y_i) \times \log(1 - p_i)\} - \frac{\vec{\beta}^T \vec{\beta}}{2 \times 10}$$

$$\text{logit}(p_i) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3$$

Since the posterior distribution is not subject to any known distribution, we need to employ the Metropolis-Hasting sampling algorithm.

Sampling Algorithm

Generally speaking, our sampling algorithm consists of the following steps:

1. Set a starting point to β
2. Sample the proposal distribution, which is centered and symmetric around $\beta^{(s)}$

$$\beta^* \sim N(\beta^{(s)}, \sigma_p^2)$$

3. Compute the acceptance ratio:

$$r = \frac{f(\beta^*|X, Y)}{f(\beta^{(s)}|X, Y)} * \frac{N(\beta^{(s)}|\beta^*)}{N(\beta^*|\beta^{(s)})}$$

4. Update the sequence:

$$\beta^{(s+1)} = \beta^* \text{ if } r > u \sim U(0, 1)$$

$$\beta^{(s+1)} = \beta^{(s)} \text{ if } r < u \sim U(0, 1)$$

5. Repeat Step2 to Step4 until a converged chain achieved

Computation Method

Let's see the sample code in Cpp, Stan, and R. Their run time will be compared later. (Stable Markets 2018)

1. Complete all sampling in R, denoted as R:

We define a function to compute the logarithm of posterior in R at very beginning.

```
# compute the log posterior of beta vector
log_posterior<-function(beta, X, Y){

  # calculate likelihood
  p_i <- invlogit(X %*% beta)
  likelihood <- sum(dbern(Y, p_i, log = T))

  # calculate prior
```

```

prior <- dmvn(x = beta, mu = rep(0,p), Sigma = sqrt((10^2)*diag(p)), log = T)

# calculate the posterior
log_post <- likelihood + prior

return(log_post)
}

```

Then a for loop is employed to generate the sample distribution with for loop.

```

# Generate the sample distribution

for(i in 2:iteration){
  beta_s <- BETA[i-1, ]

  # sample from proposal distribution
  beta_star <- mvnrm(n = 1, beta_s, Sigma = sqrt(jump_v*diag(p)))

  # calculate acceptance probability
  r_num <- log_posterior(beta_star, X, Y ) +
    dmvmnorm(x = beta_s, mean = beta_star, sigma = sqrt(jump_v*diag(p)))
  r_denom <- log_posterior(beta_s, X, Y ) +
    dmvmnorm(x = beta_star, mean = beta_s, sigma = sqrt(jump_v*diag(p)))

  r <- r_num - r_denom
  rmin<-min(r,log(1))

  # accept or reject proposal
  if( rmin > log(runif(n = 1, min = 0, max = 1)) ){
    BETA[i, ] <- beta_star
  }else{
    BETA[i, ] <- beta_s
  }
  accept[i] <- rmin
}

```

2. Compute the posterior in C++ and simulate the sample distribution in R, denoted as RCpp:

We define a function to compute the logarithm of posterior in C++ at very beginning.

```

// [[Rcpp::depends(RcppArmadillo)]]
#include <RcppArmadillo.h>
#include <Rcpp.h>
#include <math.h>
using namespace Rcpp;
using namespace arma;

vec invlogit(vec x){
  vec y = exp(x)/(1 + exp(x));
  return(y);
}

```

```

// [[Rcpp::export]]
double log_post(vec b, vec Y, mat X, vec b_s, double jump_v) {
  int n = X.n_rows;
  int p = X.n_cols;
  vec xb;
  vec pi;
  vec mu;
  mat sd_prior(p, p);
  mat sd(p, p);
  double likelihood = 0;
  double prior = 0;
  double proposal;
  double log_posterior = 0;

  // linear predictor
  xb = X*b;

  // correct for numerical issues
  for( int i=0; i<n; i++){
    if( xb[i] > 10 ){
      xb[i] = 10;
    }else if( xb[i] < -10 ){
      xb[i] = -10;
    }else{
      xb[i] = xb[i];
    }
  }

  // apply inverse link function
  pi = invlogit(xb);

  // compute log likelihood
  for(int i=0; i<n; i++){
    likelihood = likelihood + R::dbinom(Y[i], 1, pi[i], 1);
  }

  // compute log prior contribution for each parameter
  for (int i = 0; i < p; i++)
  {
    prior = prior + R::dnorm4(b[i], 0, 10, 1);
  }

  // compute the proposal density
  for (int i = 0; i < p; i++)
  {
    proposal = proposal + R::dnorm4(b_s[i], b[i], sqrt(jump_v), 1);
  }

  // evaluate log posterior as sum of likelihood and prior
  log_posterior = likelihood + prior;
  return log_posterior;
}

```

Then by using the `sourceCpp` function, we can import this C++ function in R. Then generate the sample

distribution by for loop in R.

3. Sampling in C++, denoted as C++

Since the for loop in R is somehow a bottleneck for computation efficiency, we would like to complete all the sampling task in C++. So we define a function in C++ to update the sequence of posterior, which is also involved the for in C++.

```
// [[Rcpp::export]]
List for_log_post(mat X, vec Y, int iteration, double jump_v){
  // create shell
  int p = X.n_cols;
  mat BETA(iteration, p);
  vec accept_rate(iteration);

  // starting value
  BETA.row(0) = ones<vec>(p).t();
  vec beta_0(p);
  mat sd(p, p);
  vec beta_c(p);
  double r_denum = 0;
  double r_num = 0;
  double r = 0;
  double rmin = 0;
  double uni = 0;

  for(int i = 1; i < iteration; i++){
    beta_0 = trans(BETA.row(i-1));

    sd = diagmat(jump_v * ones<vec>(p));

    // sample from proposal distribution
    beta_c = mvnrnd(beta_0, sd);

    // ratio computation
    r_num = log_post(beta_c, Y, X, beta_0, jump_v);

    r_denum = log_post(beta_0, Y, X, beta_c, jump_v);

    // acceptance ratio
    r = exp(r_num - r_denum);
    rmin = fmin(r, 1);

    //decision
    uni = randu<double>();
    if (rmin > uni){
      BETA.row(i) = beta_c.t();
    } else {
      BETA.row(i) = beta_0.t();
    }
    accept_rate(i) = rmin;
  }
  List result;
  result["BETA"] = BETA;
```

```

    result["accept_rate"] = accept_rate;
    return result;
}

```

4. RStan We define the Bayesian model in Stan as following:

```

data {
  int N;
  int p;
  matrix[N, p] x;
  int y[N];
}

parameters {
  vector[p] beta;
}

model{
  // prior
  beta ~ normal(0, 100);

  //likelihood
  y ~ bernoulli_logit(x * beta);
}

```

Then we can implement this model by commanding in R:

```

res_rstan <- stan(file = 'posterior_in_stan.stan',
                 data = list(x = X, y = Y, N = nrow(X), p = ncol(Y)),
                 iter = iteration, chains = 1)

```


Run Time Comparison

Run Time by Method

All these approaches will generate similar posterior distribution. I attach the trace plots for these different as appendix. Now let's focus on the run time of these different approaches. The R function `microbenchmark` is used here to measure the running time accurately, which is a good substitution for the `system.time(replicate(repeat_time, function))` expression. The result is as following:

Table 1: Run Time Summary with Sample Size 50

| method | min | lq | mean | median | uq | max | neval |
|--------|------------|------------|------------|------------|-----------|------------|-------|
| Pure R | 16541.4761 | 17448.0089 | 18159.5563 | 17805.9182 | 18646.854 | 22250.6081 | 100 |
| RCpp | 1673.0167 | 1757.6993 | 1840.5268 | 1804.0076 | 1899.491 | 2356.4302 | 100 |
| C++ | 251.3411 | 258.2004 | 290.7577 | 273.3687 | 299.332 | 555.1278 | 100 |
| Stan | 32026.7448 | 33644.3007 | 34746.3294 | 34243.7151 | 35054.643 | 40277.1960 | 100 |



Conclusion

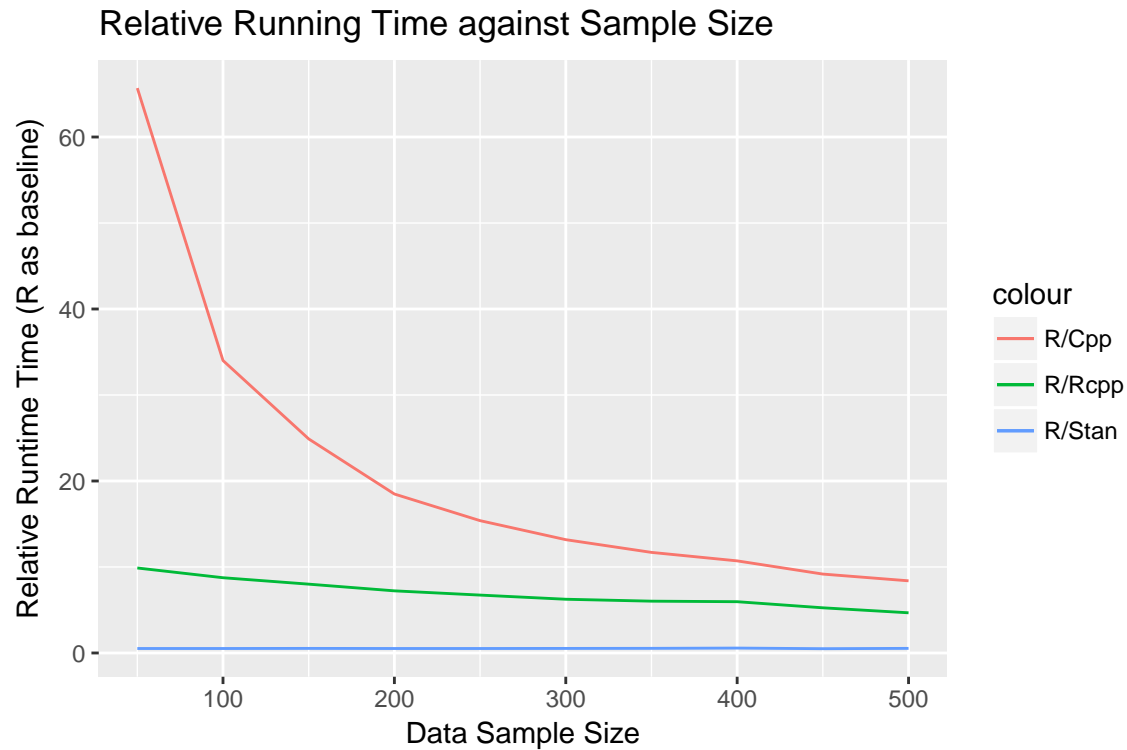
As we can see, if we generate the sample distribution in C++, the run time can be much shorter than any other method. Even when we only put the computation of posterior in C++, it can still save around 90% run time compared to the pure R computation. However, due to time cost in communication between R and Stan, RStan returns the worst run times. If we also consider the hours spent on coding and mathematics deduction, RStan will be the most efficient one.

Run Time by Sample Size

In previous case, we study the case with sample size 50. We are also interested in the relationship between the run time and the sample size, which is from 50 to 500 with increment of 50. We set the run time of R as the baseline.

Table 2: Relative Run Time Summary against Sample Size

| Sample Size | R/Rcpp | R/Cpp | R/Stan |
|-------------|----------|-----------|-----------|
| 50 | 9.882080 | 65.689113 | 0.5193266 |
| 100 | 8.749257 | 34.008443 | 0.5206681 |
| 150 | 8.003832 | 24.891335 | 0.5331680 |
| 200 | 7.225733 | 18.484568 | 0.5202131 |
| 250 | 6.730718 | 15.387929 | 0.5200740 |
| 300 | 6.248238 | 13.175251 | 0.5285307 |
| 350 | 6.023071 | 11.696660 | 0.5369588 |
| 400 | 5.961274 | 10.709417 | 0.5688523 |
| 450 | 5.253168 | 9.174063 | 0.4965358 |
| 500 | 4.673372 | 8.391021 | 0.5362079 |

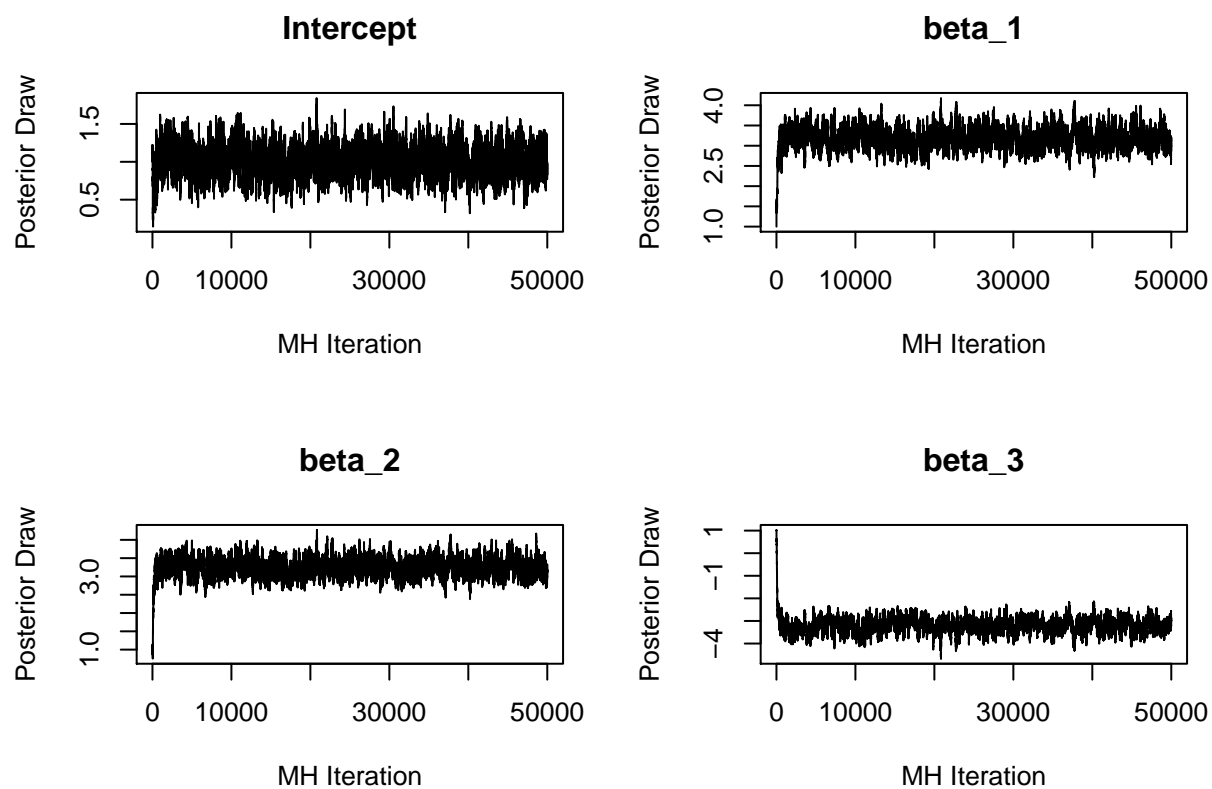


Conclusion

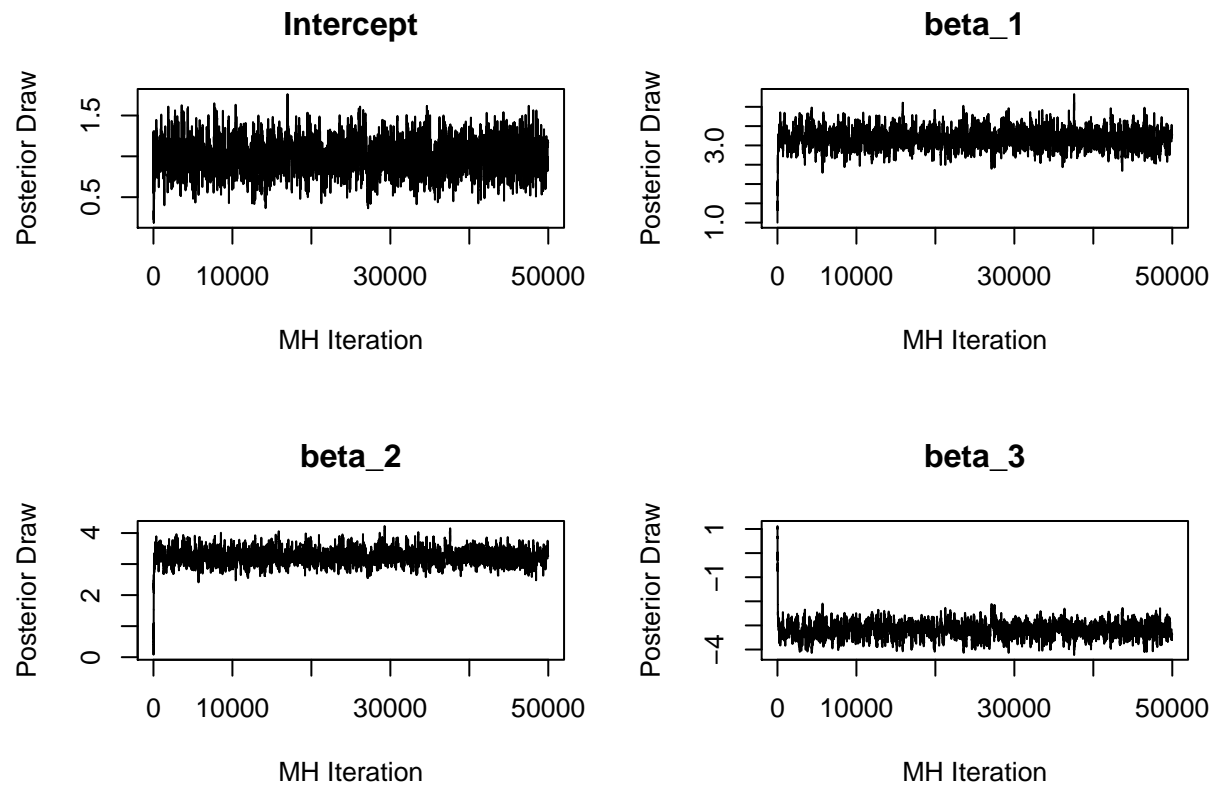
As we can see, C++ method is always dominating. As the sample size increase, all these run time ratios are dropping. And an interesting result is from Stan. Although built on C++, Stan is even slower than R method. We believe the tuning procedure in Stan costs significantly in the 'warm-up' procedure, which will automatically find out the best proposal distribution. Therefore, if you want to speed up your sampling, you can try to sample in C++.

Appendix: Trace Plot for Different Methods

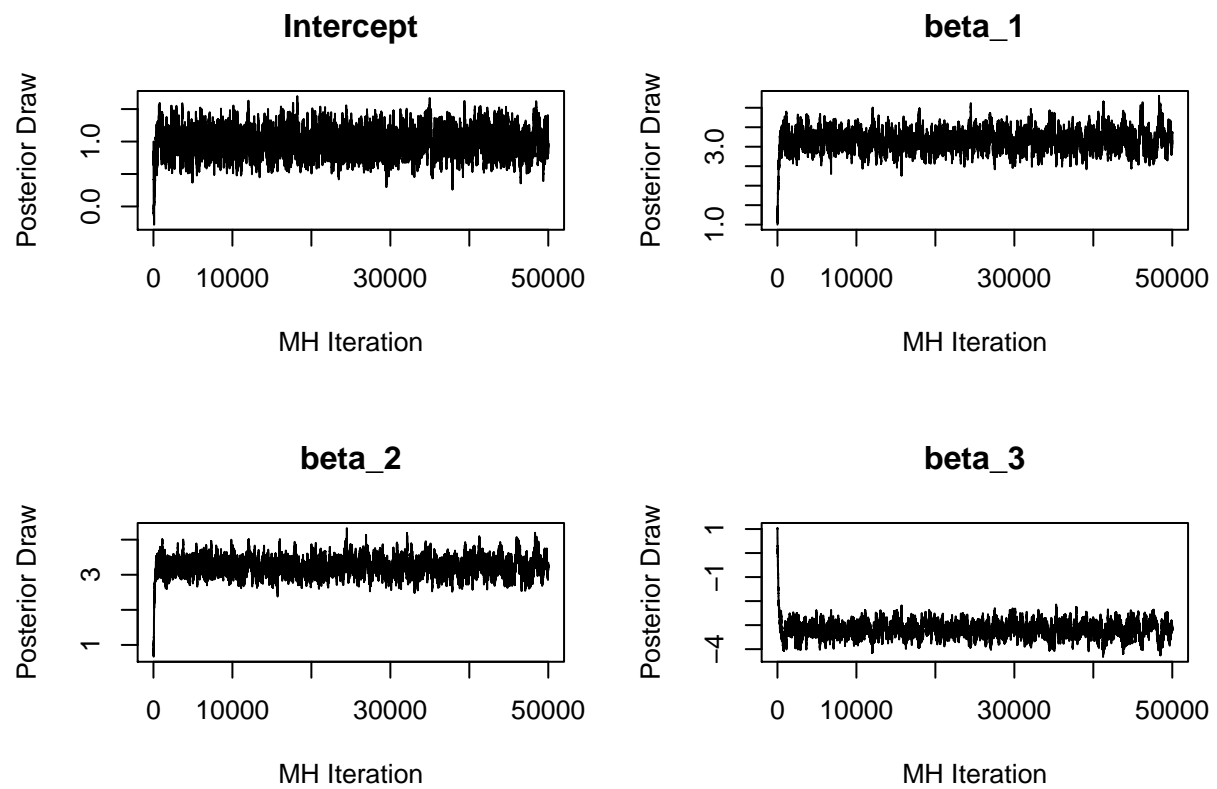
1. R



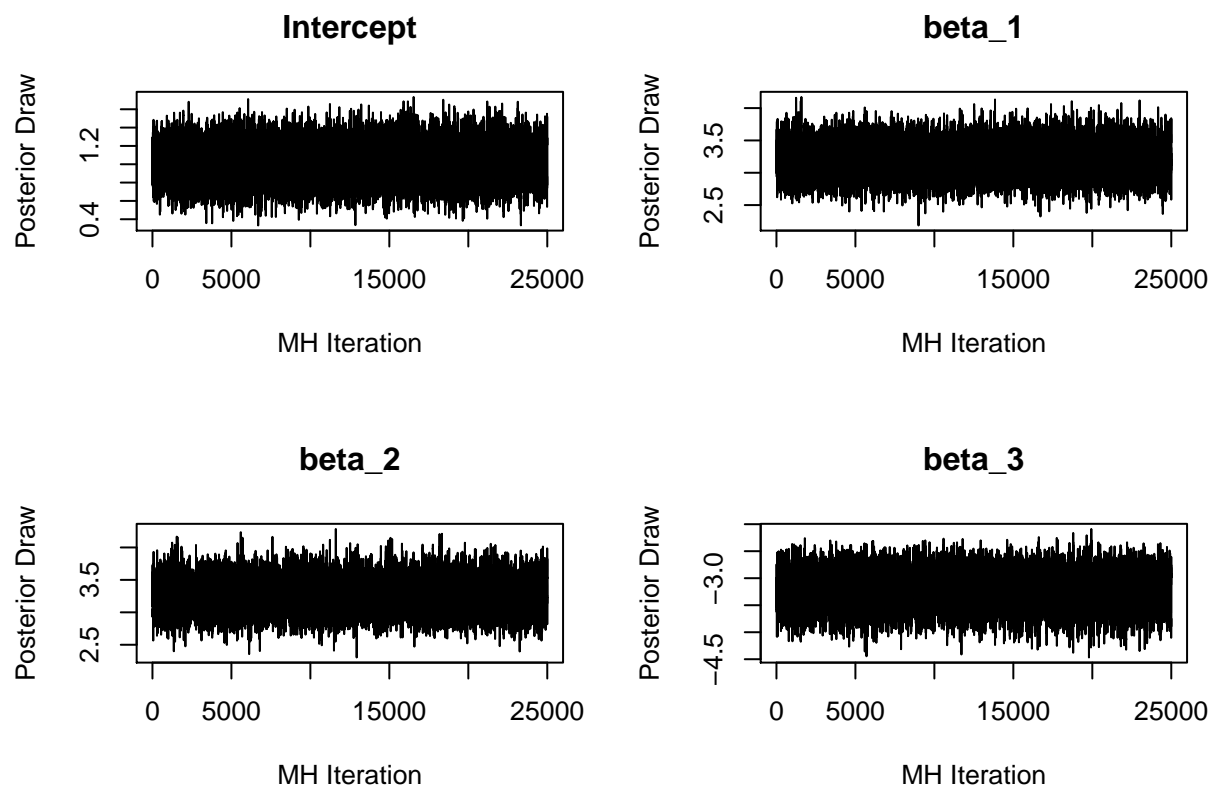
2. RCpp



3. Cpp



4. RStan



Reference

Eddelbuettel, Dirk, Romain François, J Allaire, Kevin Ushey, Qiang Kou, N Russel, John Chambers, and D Bates. 2011. “Rcpp: Seamless R and C++ Integration.” *Journal of Statistical Software* 40 (8): 1–18.

RCpp. 2018. “Rcpp for Seamless R and C++ Integration.” <http://www.rcpp.org/>.

Stable Markets. 2018. “Speeding up Metropolis-Hastings with Rcpp.” <https://stablemarkets.wordpress.com/2018/03/16/speeding-up-metropolis-hastings-with-rcpp/>.

Wikipedia contributors. 2018a. “C++ — Wikipedia, the Free Encyclopedia.” <https://en.wikipedia.org/wiki/C%2B%2B>.

———. 2018b. “MCMC — Wikipedia, the Free Encyclopedia.” https://en.wikipedia.org/wiki/Markov_chain_Monte_Carlo.